

NOT A REPLACEMENT FOR THE GUIDELINES. Section refs point into *[External] OpenClaw MM Rubrics MULTI TURN – Guidelines*, which stays the source of truth. Tick a box only when you have actually looked, not when you assume.

1 Scenario shape

- ☐ **A1** Does the scenario sit naturally inside the assigned category, subcategory and universe, rather than having been bent to fit them?
If not, re-read the subcategory definition and adjust until the pair is the scenario's obvious home. [§1.1](#) · [1.2.5](#)
- ☐ **A2** Did you ground the scenario by interacting with the AI agent in the **Database tab**, confirm the loadout in the Universe Explorer, add the Service Universe Artifact ID, and anchor every date to a window you actually saw in the data?
Inspecting by hand is what leads people to conclude the scenario has no anchor in the universe. Without the Artifact ID the artifacts come back empty, and an unloaded server is an environment defect, not a model failure. [§1.1](#) · [1.2.1](#)
- ☐ **A3** Does the task require multi-system coordination across all three stages — data acquisition, processing and reasoning, output generation?
If it can be completed in a short linear exchange it is too simple. Add a cross-system handoff, not more asks. [§Hard Client Requirements](#) · [1.2.5](#)
- ☐ **A4** Did Model A actually fail at least 30% of the final rubric score, with failures that materially affect task completion?
If the run sailed through, raise the reasoning difficulty. Do not close the gap with format micro-specifications the model happens to miss. [§Hard Client Requirements](#)

2 Universe and evidence

- ☐ **B1** Is there at least one fact the task cannot be completed without that lives **only** in a connected service, and does the agent require meaningful interaction with the universe rather than superficial?
- ☐ **B2** Take the attachments away — does the task become unsolvable? If no, the model does not require the level of multimodal reasoning expected.
Not every input needs to be essential: realistic noise and distractor files are allowed and encouraged, as long as the task still genuinely requires multimodal reasoning to recover the facts needed for completion. [§1.2.2](#)
- ☐ **B3** Do the inputs clear the hard rules — downsampled, no `.heic`, no LLM-generated `.pdf/ .docx/ .xlsx`, under the 20% LLM cap, no junk or system files, CCO or CC BY, synthetic personas only, and **at least three of them** in the conversation?
Three input files is the floor; distractors and noise count toward it. Then check the names: no filename, manifest or helper document may reveal the expected answer. [§1.2.2](#)

3 Prompt and turns

- ☐ **C1** Is every expected output file named explicitly in the prompt itself, spelled exactly as it must appear?
A filename that appears only in an input file, the Agent Objective or the Desired Outcome was never asked for. [§1.2.3](#) · [1.3](#)
- ☐ **C2** Do the thresholds, rules and policies the agent must apply live anywhere in the scenario (prompt, multimodal inputs)?
[§1.2.3](#)
- ☐ **C3** Does the conversation run 3 to 5 turns, with every follow-up consuming the state the turn before it produced?
Fewer than 3 turns is automatically rejected. A follow-up that would work as an independent opening prompt fails Turn Structure. [§1.1](#) · [1.2.3](#)
- ☐ **C4** Is there at least one turn that changes the brief after something was already delivered, and is any deferred asset named explicitly on the turn it enters scope?
A revision is a targeted edit to what exists. "I found another file" does not name the asset. [§Hard Client Requirements](#) · [1.2.3](#) · [6.3](#)
- ☐ **C5** Did you avoid flagging the miss anywhere — in the Leg A follow-ups and in your Leg B both follow ups and hints?
Pointing at the failure compromises the 30% threshold. If the model could copy your hint straight into the artifact, you gave away the answer. [§1.2.3](#) · [7.2](#)

4 Draft History alignment

- ☐ **D1** Is every requirement you intend to grade actually stated in a prompt the agent received?
The Draft History is not embedded in the agent. A requirement living only there was never asked for, so it cannot be graded. [§1.3](#) · [4](#)
- ☐ **D2** Does the Desired Outcome describe the shape of the end state without pre-filling the answers, and does it match the GTFA you resolved?
If you cannot state the one correct answer yourself, the scenario was never ready to write up. [§1.2.4](#) · [1.3](#)

☐ Ready to submit

Every box ticked, and each one because you looked rather than assumed. If a check felt uncertain rather than clean, it is not ticked — go back to the section it names.

5 Objective rubrics

- ☐ **E1** Walking the prompt once per turn, is every ask covered by a criterion — including intents introduced in the middle turns, not just the final state?
Missing Criteria and Turn Scoped are both Major issues. No target count: thorough coverage usually lands around 15 to 30. [§5.1](#) · [5.6](#) · [5.7](#)
- ☐ **E2** Read with the prompt closed: can a grader rate every criterion Present or Not Present from the criterion and its evaluation target alone?
Embed the exact value, filename, date or classification. Anchor "the photo" to the actual filename. [§5.1](#)
- ☐ **E3** Is every value asserted in a criterion the one the tool, file or universe actually returns at runtime?
Verify against the downloaded trajectory and artifacts rather than the note you took while authoring. [§5.1](#) · [5.5](#)
- ☐ **E4** Does every literal format requirement trace back to a line you can quote from the prompt?
Exact heading names, ordering, casing and phrasing are Instruction Following only when the user stated them. If you cannot cite the line, cut it. [§5.1](#) · [5.4](#)
- ☐ **E5** Is each criterion atomic, and do no two criteria demand mutually exclusive outcomes or check the same thing twice with opposite polarity?
Split only where two halves could pass independently; several values serving one finding stay in one criterion. A positive on X and a negative on Y are complementary, not contradictory. [§5.1](#) · [5.6](#)
- ☐ **E6** Is every weight in $\{-5, -3, -1, +1, +3, +5\}$, and does each answer how hard the criterion was to satisfy rather than how much it matters?
Any value outside the set fails Rubric Structure on its own. A requirement can be critical to the task and still be +1. [§5.2](#) · [5.7](#)
- ☐ **E7** Is every criterion in the category and evaluation target it would actually be graded under?
Agent Behavior is always Trajectory — if you wrote it against an artifact, a state change or the final message, you are grading the deliverable and the category is wrong. **Factuality & Hallucination has left the table:** a criterion checking whether a value or claim is correct is Task Completion. [§5.4](#)
- ☐ **E8** Are negatives around a quarter of the block and under 30%, and does any group of more than eight similar outcomes use one completeness criterion plus at most five spot checks?
Each negative must name a failure the setup genuinely invites, not mirror every "don't" in the prompt. [§5.1.1](#) · [5.3](#)
- ☐ **E9** Does the block carry **five or fewer** criteria whose evaluation target is Trajectory?
Five is the ceiling and zero is allowed; they are not mandatory, and they focus on the last response. Over five, move the weakest onto the artifacts rather than dropping the coverage. [§5.1](#)

6 Subjective block

- ☐ **F1** Does the block hold 10 or more natural criteria, each anchored to one identifiable element and one visible property the prompt never asked for?
Do not force extra rubrics to hit the number. No "looks professional", "well designed" or "high quality" — name the property a reviewer can locate and score. [§8.1](#)
- ☐ **F2** Do the weights measure impact on the user's experience rather than difficulty, and does every criterion grade something the format can actually show?
A PDF cannot respond to hover and a static graphic cannot play audio. Task Completion and Final Answer Artifact are the only valid category and target. [§8.1](#) · [8.2](#) · [8.4](#)
- ☐ **F3** Did you rate every subjective criterion Present or Not Present against **both** models yourself, and are the justifications written at the user-experience level?
Pre-filled selections are not to be trusted, and the Model B set is entirely yours to determine. [§8.5](#)

7 Dynamic values and final state

- ☐ **G1** If the scenario touches the public internet or other dynamic type of information (essentially prices), is no criterion pinned to a value that can legitimately change between runs?
Grade the retrieval and the relationship the model has to establish, with the value given as an example rather than as the required answer. [§9.3.1](#)
- ☐ **G2** Does the golden pass the complete objective rubric set, with every named file present and every factual claim grounded in an input that was actually opened?
Anything the golden fails is a broken criterion, not a broken golden. [§5.7](#) · [7](#)
- ☐ **G3** Are the trajectories and final artifacts downloaded, the preferred run starred, and does every failed criterion carry all three justification parts?
A failure is a positive rated Not Present or a negative rated Present. Every one carries a justification; the golden needs none. [§5.5](#) · [7.1](#)

☐ Fix before submitting

Any box left empty. Fix the item, then re-run the section it sits in — changing a rule in the prompt or the inputs usually moves something in the rubric block too.