

**NOT A REPLACEMENT FOR THE GUIDELINES.**

Section refs point into *[External] OpenClaw MM Rubrics MULTI TURN – Guidelines*, which stays the source of truth. Tick a box only when you have actually looked, not when you assume.

## 1 Parameters and scenario shape

*Is this the task you were assigned, and is it hard enough to be worth grading?*

**Four parameters are assigned and cannot change**

Task type, category, subcategory, universe. Drift from any one of them is rejected at QC, however good the task is.

**Universe dates are static**

✗ “next Tuesday” — can land on a week that holds nothing

✓ “the week of 12 May” — a range you actually saw

- ☐ **A1** Does the scenario sit naturally inside the assigned category, subcategory and universe, rather than having been bent to fit them?  
If not, re-read the subcategory definition and adjust until the pair is the scenario's obvious home. §1.1 · 1.2.5
- ☐ **A2** Did you confirm the loadout in the Universe Explorer, add the Service Universe Artifact ID, and anchor every date to a window you actually saw in the data?  
Without the Artifact ID the artifacts come back empty. Designing against a server that is not loaded is an environment defect, not a model failure. §1.1 · 1.2.1
- ☐ **A3** Does the task genuinely need three stages and two or more systems, with one real friction point?  
If it can be completed in a short linear exchange it is too simple. Add a cross-system handoff, not more asks. §Hard Client Requirements
- ☐ **A4** Did Model A actually fail at least 50% of the final rubric score, with failures that materially affect task completion?  
If the run sailed through, the scenario needs more difficulty. Do not close the gap by adding criteria the model happens to miss. §Hard Client Requirements

## 2 Universe and evidence

*Is the media load-bearing, and is the universe doing real work?*

**Four ways this goes wrong**

**Write only** — every figure comes from the attachments, the universe only written to.

**Universe dropped** — the media settles everything.

**Pictures of the universe** — the input duplicates what the universe holds. **A cross-check nobody could perform.**

**Observed in a real run**

282 service tools loaded across 14 servers. The agent called **none** of them, even with the prompt naming the Slack channel twice. Connected does not mean used.

**Messy is required, illegible is a defect**

✓ Blurry phone shots, skewed scans, IMG\_0427.jpg

✗ A value no person could read → feasibility failure

- ☐ **B1** Is there at least one fact the task cannot be completed without that lives **only** in a connected service, and does the agent have to read the universe rather than only write to it?  
Notes, drafts and calendar entries the agent creates are outputs, not universe interaction. §1.2.2 · 9.2
- ☐ **B2** Take the attachments away — does the task become unsolvable?  
If any required value also exists as text somewhere in the environment, MM dependence fails and the media reads as decorative. §1.2.2
- ☐ **B3** Does every attached file carry at least one needed fact, and could a person read every load-bearing value at full size?  
A file that can be left unopened is junk and can get the task rejected. §1.2.2 · 4
- ☐ **B4** Do the inputs clear the hard rules — downsampled, no .heic, no LLM-generated .pdf/.docx/.xlsx, under the 20% LLM cap, no junk or system files, CC0 or CC BY, synthetic personas only?  
Then check the names: no filename, manifest or helper doc may reveal an answer or signal which turn a deferred asset belongs to. §1.2.2

## 3 Prompt and turns

*Does the conversation actually ask for what you are about to grade?*

**Name every output file, in the prompt**

✗ “generate a report with your findings”

✓ “write your findings in final\_report.md”

**3 to 5 turns**

Opening prompt plus 2 to 4 follow-ups. Fewer than 3 is automatically rejected. Each follow-up must consume the state the one before it produced.

**Never flag the miss**

Not in the Leg A follow-ups, not in the Leg B hints. The test: if the model could copy your hint straight into the artifact, you gave away the answer.

- ☐ **C1** Is every expected output file named explicitly in the prompt itself, spelled exactly as it must appear?  
A filename that appears only in an input file, the Agent Objective or the Desired Outcome was never asked for. §1.2.3 · 1.3
- ☐ **C2** Do the thresholds, rules and policies the agent must apply live in the input files rather than being stated in the prompt?  
Stating the policy in the prompt removes the work you intended to grade. §1.2.3
- ☐ **C3** Does the conversation run 3 to 5 turns, with every follow-up consuming the state the turn before it produced?  
A follow-up that would work as an independent opening prompt fails Turn Structure. §1.1 · 1.2.3 · 3
- ☐ **C4** Is there at least one turn that changes the brief after something was already delivered, and is any deferred asset named explicitly on the turn it enters scope?  
A revision is a targeted edit to what exists. “I found another file” does not name the asset. §Hard Client Requirements · 6.3
- ☐ **C5** Did you avoid flagging the miss anywhere — in the Leg A follow-ups and in your Leg B hints?  
Pointing at the failure compromises the 50% threshold. §1.2.3 · 7.2

## 4 Draft History alignment

*Nothing the agent never saw can be graded.*

**The agent sees only your prompts and attachments**

The Draft History, Agent Objective, Desired Outcome, rubrics and milestones are invisible to it — the most common invalid assumption in task design. The Desired Outcome names each artifact, its contents and the decision logic, and never the values the agent has to derive.

- ☐ **D1** Is every requirement you intend to grade actually stated in a prompt the agent received?  
A requirement living only in the Draft History was never asked for, so it cannot be graded. §1.3 · 4
- ☐ **D2** Does the Desired Outcome describe the shape of the end state without pre-filling the answers, and does it match the GTFA you resolved?  
If you cannot state the one correct answer yourself, the scenario was never ready to write up. §1.2.4 · 1.3

## 5 Objective rubrics

Every criterion has to be ratable, correct at runtime, and grounded in something the prompt said.

### Severity thresholds — any one of these fails the task

**MAJOR** above 10% of the set · **MODERATE** above 15% · **MINOR** above 20%. Two to four issues is often enough.

### Self-containment

- ☒ maintenance\_summary.xlsx contains the correct inspection date for unit 7C
- ☒ maintenance\_summary.xlsx records 2026-03-14 as the inspection date for unit 7C

### Values that drift by runtime

- ☒ \$25.00 in the rubric, \$35.00 from the tool
- ☒ “Ikenna” in the rubric, “Ifeanyi” from the calendar
- ☒ 1080p in the rubric, 1920x1080 from the file  
These are rubric defects, not model failures.

### Format over-specification

- ☒ “exactly 11 ## headings in this order” when the prompt said “include these sections”
- ☒ “heading must be ## Summary” when the prompt said “wrap up with a summary”

### Two rules people get wrong

Weight is **difficulty**, **never importance** — a criterion can be critical and still be +1. Agent Behavior is **always Trajectory**.

- ☐ **E1** Walking the prompt once per turn, is every ask covered by a criterion — including intents introduced in the middle turns, not just the final state?  
Missing Criteria and Turn Scoped are both Major issues. §5.1 · 5.6 · 5.7
- ☐ **E2** Read with the prompt closed: can a grader rate every criterion Present or Not Present from the criterion and its evaluation target alone?  
Embed the exact value, filename, date or classification. Anchor “the photo” to the actual filename. §5.1
- ☐ **E3** Is every value asserted in a criterion the one the tool, file or universe actually returns at runtime?  
Verify by making the call yourself rather than trusting the note you took while authoring. §5.1 · 5.5
- ☐ **E4** Does every literal format requirement trace back to a line you can quote from the prompt?  
Exact heading names, ordering, casing and phrasing are Instruction Following only when the user stated them. If you cannot cite the line, cut it. §5.1 · 5.4
- ☐ **E5** Is each criterion atomic, and do no two criteria demand mutually exclusive outcomes or check the same thing twice with opposite polarity?  
Three entities means three criteria. Check polarity first: a positive on X and a negative on Y are complementary, not contradictory. §5.1 · 5.6
- ☐ **E6** Is every weight in {−5, −3, −1, +1, +3, +5}, and does each answer how hard the criterion was to satisfy rather than how much it matters?  
Any value outside the set fails Rubric Structure on its own. §5.2 · 5.7
- ☐ **E7** Is every criterion in the category and evaluation target it would actually be graded under?  
If you wrote Agent Behavior against an artifact, a state change or the final message, you are grading the deliverable and the category is wrong. §5.4
- ☐ **E8** Are negatives around a quarter of the block and under 30%, and does any group of more than eight similar outcomes use one completeness criterion plus at most five spot checks?  
Each negative must name a failure the setup genuinely invites, not mirror every “don’t” in the prompt. §5.1.1 · 5.3

## 6 Subjective block

Presentation only, judged on the render.

### One element, one visible property

- ☒ island\_ferry\_times.pdf has an appropriate and professional look
- ☒ In island\_ferry\_times.pdf, every departure time is aligned on its colon in a single column

### Scope

Task Completion is the only valid category, Final Answer Artifact the target. Anything about the render that the prompt asked for is **objective only**. Weights measure user-experience impact, never difficulty.

- ☐ **F1** Does the block hold 10 or more criteria, each anchored to one identifiable element and one visible property the prompt never asked for?  
No “looks professional”, “well designed” or “high quality” — name the property a reviewer can locate and score. §8.1
- ☐ **F2** Do the weights measure impact on the user’s experience rather than difficulty, and does every criterion grade something the format can actually show?  
A PDF cannot respond to hover and a static graphic cannot play audio. §8.1 · 8.2
- ☐ **F3** Did you rate every subjective criterion Present or Not Present against **both** models yourself, and are the justifications written at the user-experience level?  
Pre-filled selections are not to be trusted, and the Model B set is entirely yours to determine. §8.4 · 8.5

## 7 Dynamic values and final state

The last things that break a task after everything else is right.

### Live web values

- ☒ reports \$129.99 as StrideHub’s current price
- ☒ reports StrideHub’s current price based on information retrieved during the run (e.g. \$129.99)

### The Trajectory button

It only downloads the **most recent** trajectory. Click it once per turn, or you will not have the evidence to grade trace-based criteria.

- ☐ **G1** If the scenario touches the public internet, is no criterion pinned to a value that can change between runs — and is nothing in the task simply lookup-able online?  
Grade the retrieval and the relationship, with the value as an example. Anything answerable online gets fetched instead of worked out. §9.3.1
- ☐ **G2** Does the golden pass the complete objective rubric set, with every named file present and every factual claim grounded in an input that was actually opened?  
Anything the golden fails is a broken criterion, not a broken golden. §5.7 · 7
- ☐ **G3** Are the trajectories and final artifacts downloaded for every turn, the preferred run starred, and does every failed criterion carry all three justification parts?  
Add by hand any justification the taxonomy did not pull automatically. §3 · 5.5

### ☐ Ready to submit

All 29 boxes ticked, and each one because you looked rather than assumed. If a check felt uncertain rather than clean, it is not ticked — go back to the section it names.

### ☐ Fix before submitting

Any box left empty. Fix the item, then re-run the section it sits in — changing a rule in the prompt or the inputs usually moves something in the rubric block too.